

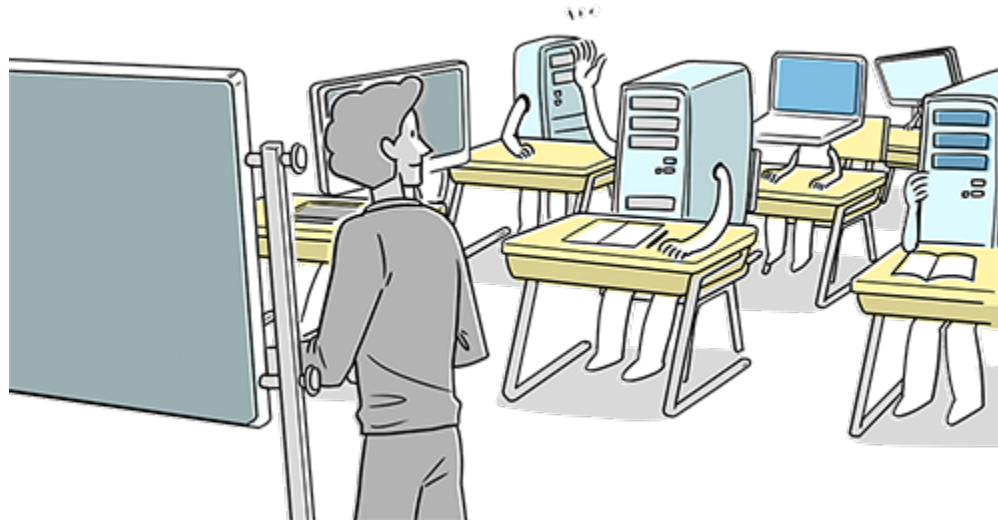


CPE/EE 695: Applied Machine Learning

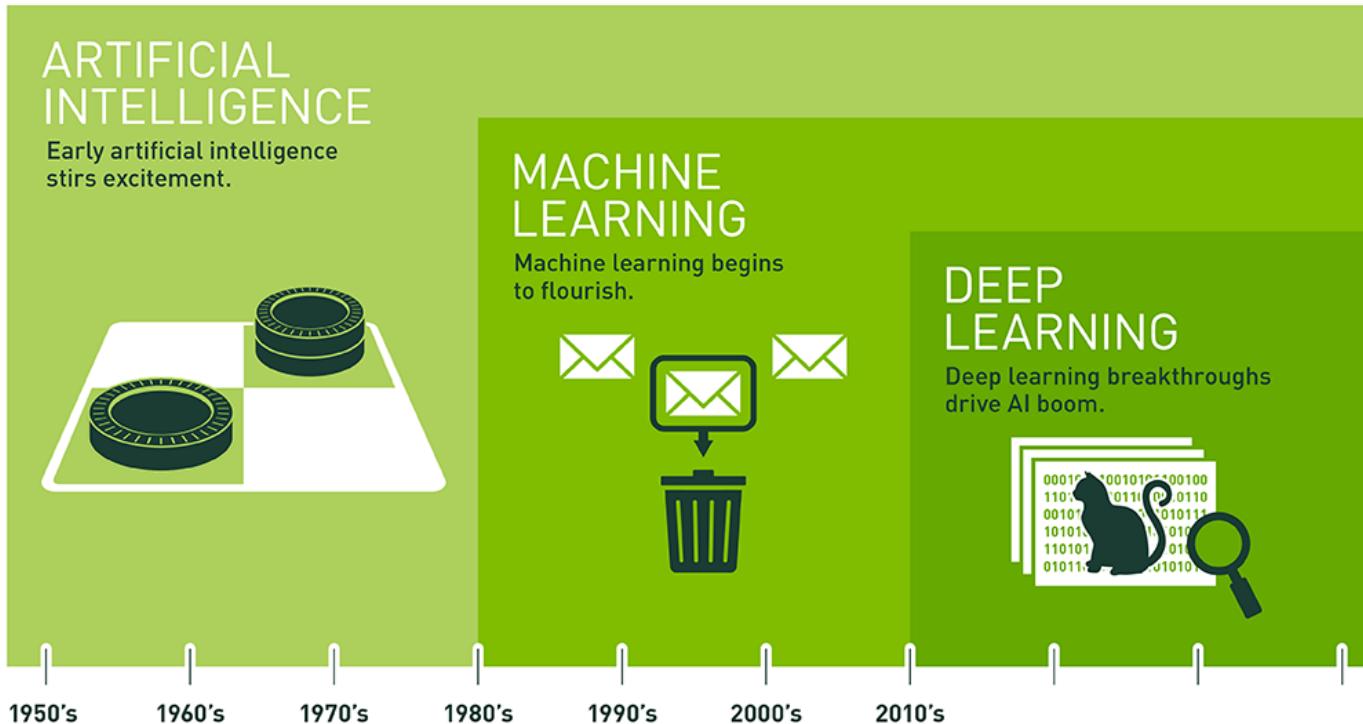
Lecture 1: Introduction to ML

Dr. Shucheng Yu, Associate Professor
Department of Electrical and Computer Engineering
Stevens Institute of Technology

What is Machine Learning?



Relation with other “buzzwords”



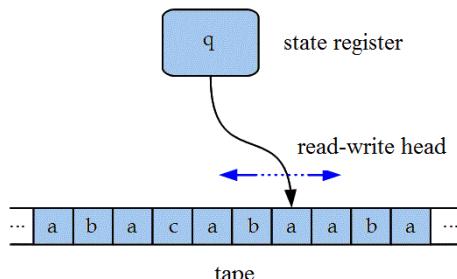
Since an early flush of optimism in the 1950s, smaller subsets of artificial intelligence – first machine learning, then deep learning, a subset of machine learning – have created ever larger disruptions.

Computability theory (1930's -)

Church-Turing Thesis (1936):

“a function on the natural numbers can be calculated by an effective method, if and only if it is computable by a **Turing Machine**”

“digital computers can simulate any process of formal reasoning”

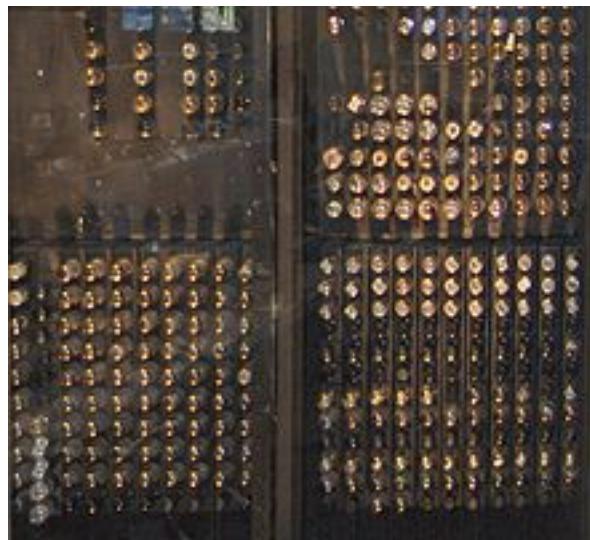


Turing Machine



The First General-Purpose Electronic Computer

ENIAC (1946) - Electronic Numeral Integrator And Computer

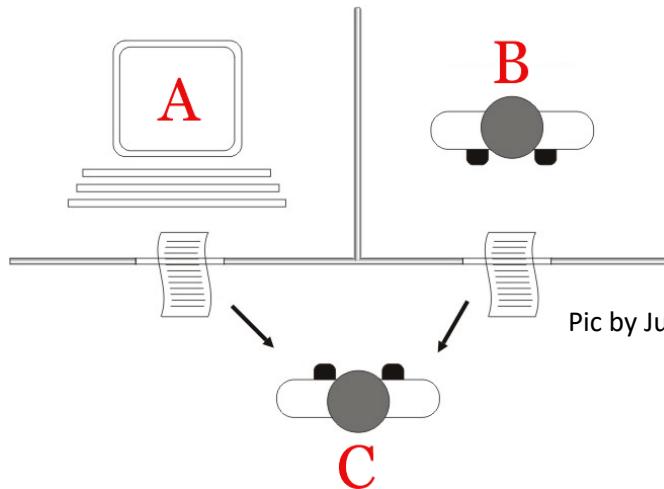


17,468 vacuum tubes,
70,000 resistors,
10,000 capacitors,
1,500 relays,
6,000 manual switches ,
5 million soldered joints

Can Machine Think?

Turing Test (1950)

- ELIZA (1966)
- PARRY(1972)
- Loebner Prize (1990)



Pic by Juan Alberto Sánchez Margallo

Player C, the interrogator, is given the task of trying to determine which player – A or B – is a computer and which is a human. The interrogator is limited to using the responses to written questions to make the determination

Can Siri pass the Turing test?



(Recorded on July 29, 2013)

Human: Hello I'm Ronan. what is your name?

Siri: Siri. Please to meet you.

Human: What is your mother's name?

Siri: I'd rather not say.

Human: What is your birth sign?

Siri: I don't really like talking about myself.

Human: How many children do you have?

Siri: Only biological entities have children.

Human: Do you prefer red or white wine?

Siri: I couldn't find any restaurants matching "red or white wine"

Human: I like bananas. Which is your favorite fruit?

Siri: I prefer not to say.

Human: What music do you like to listen to?

Siri: I really have no opinion.

Human: what is your favorite song?

Siri: I like whatever you like.

Human: I like Waiting for Godot. What is your favorite play?

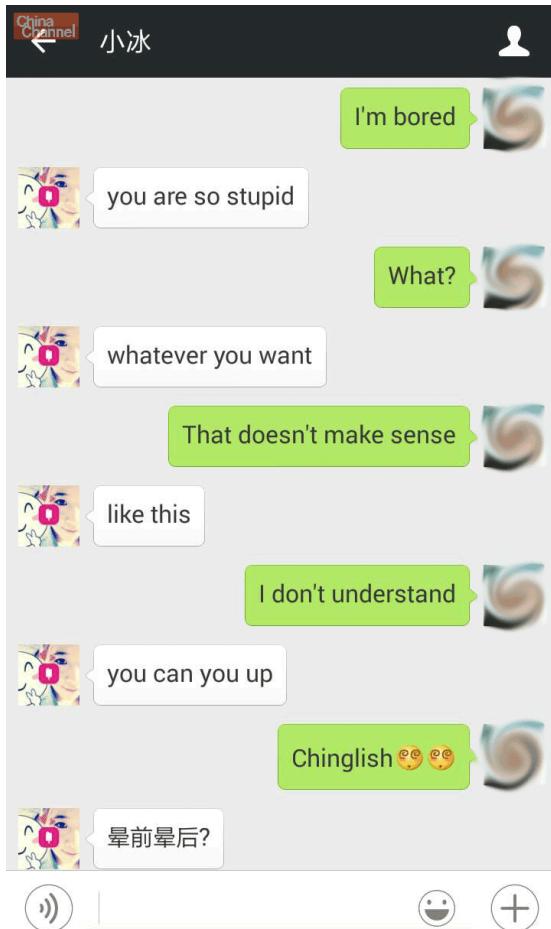
Siri: I don't really like talking about myself.

Human: What color do you dye your hair?

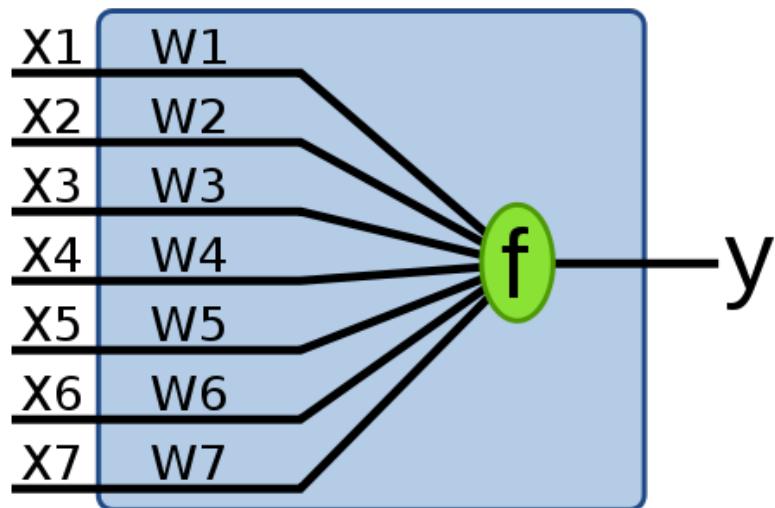
Siri: Would you like me to search the web for "What color do you dye your hair?"?



Microsoft Chatbot Xiaobing



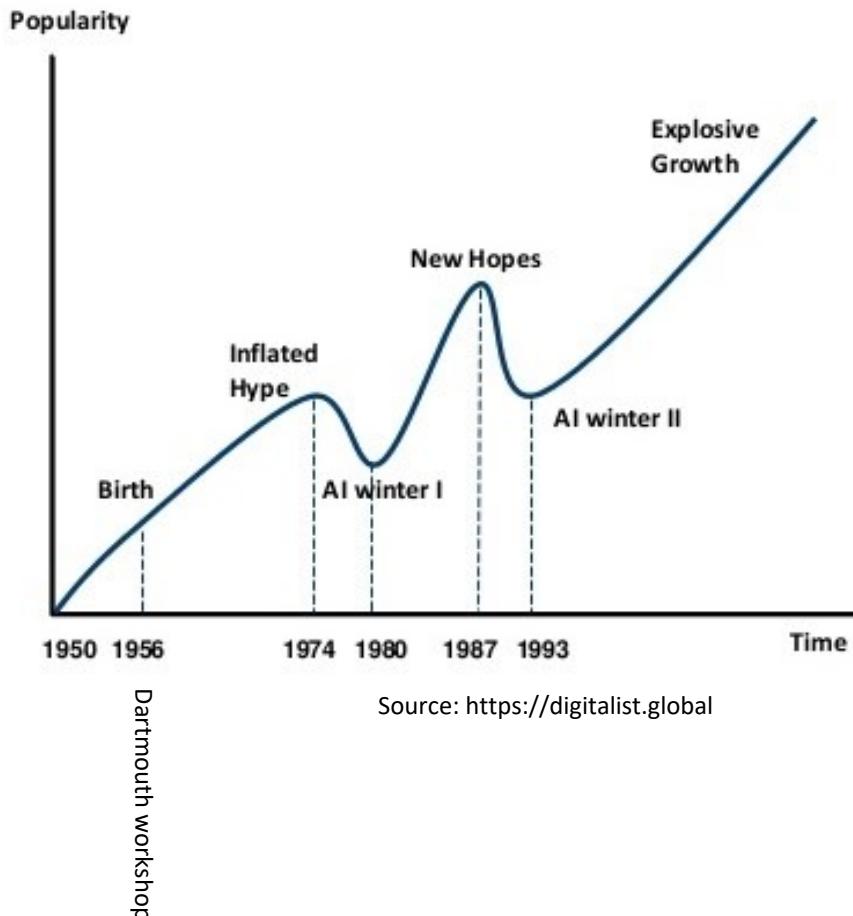
Perceptrons



$$f(x) = \begin{cases} 1, & \text{if } wx + b > 0 \\ 0, & \text{otherwise} \end{cases}$$

AI Cycles

AI HAS A LONG HISTORY OF BEING “THE NEXT BIG THING”...



Timeline of AI Development

- **1950s-1960s:** First AI boom - the age of reasoning, prototype AI developed
- **1970s:** AI winter I
- **1980s-1990s:** Second AI boom: the age of Knowledge representation (appearance of expert systems capable of reproducing human decision-making)
- **1990s:** AI winter II
- **1997:** Deep Blue beats Gary Kasparov
- **2006:** University of Toronto develops Deep Learning
- **2011:** IBM's Watson won Jeopardy
- **2016:** Go software based on Deep Learning beats world's champions

AI Winter

- Failure of Machine Translation -- disambiguation
- Perceptron Limitation – Linear
- Failure of Speech Understanding Research at CMU
- Lighthill Report(1973): James Lighthill, “*Artificial Intelligence: a paper symposium*”

"In no part of the field have the discoveries made so far produced the major impact that was then promised"

- DARPA funding cut

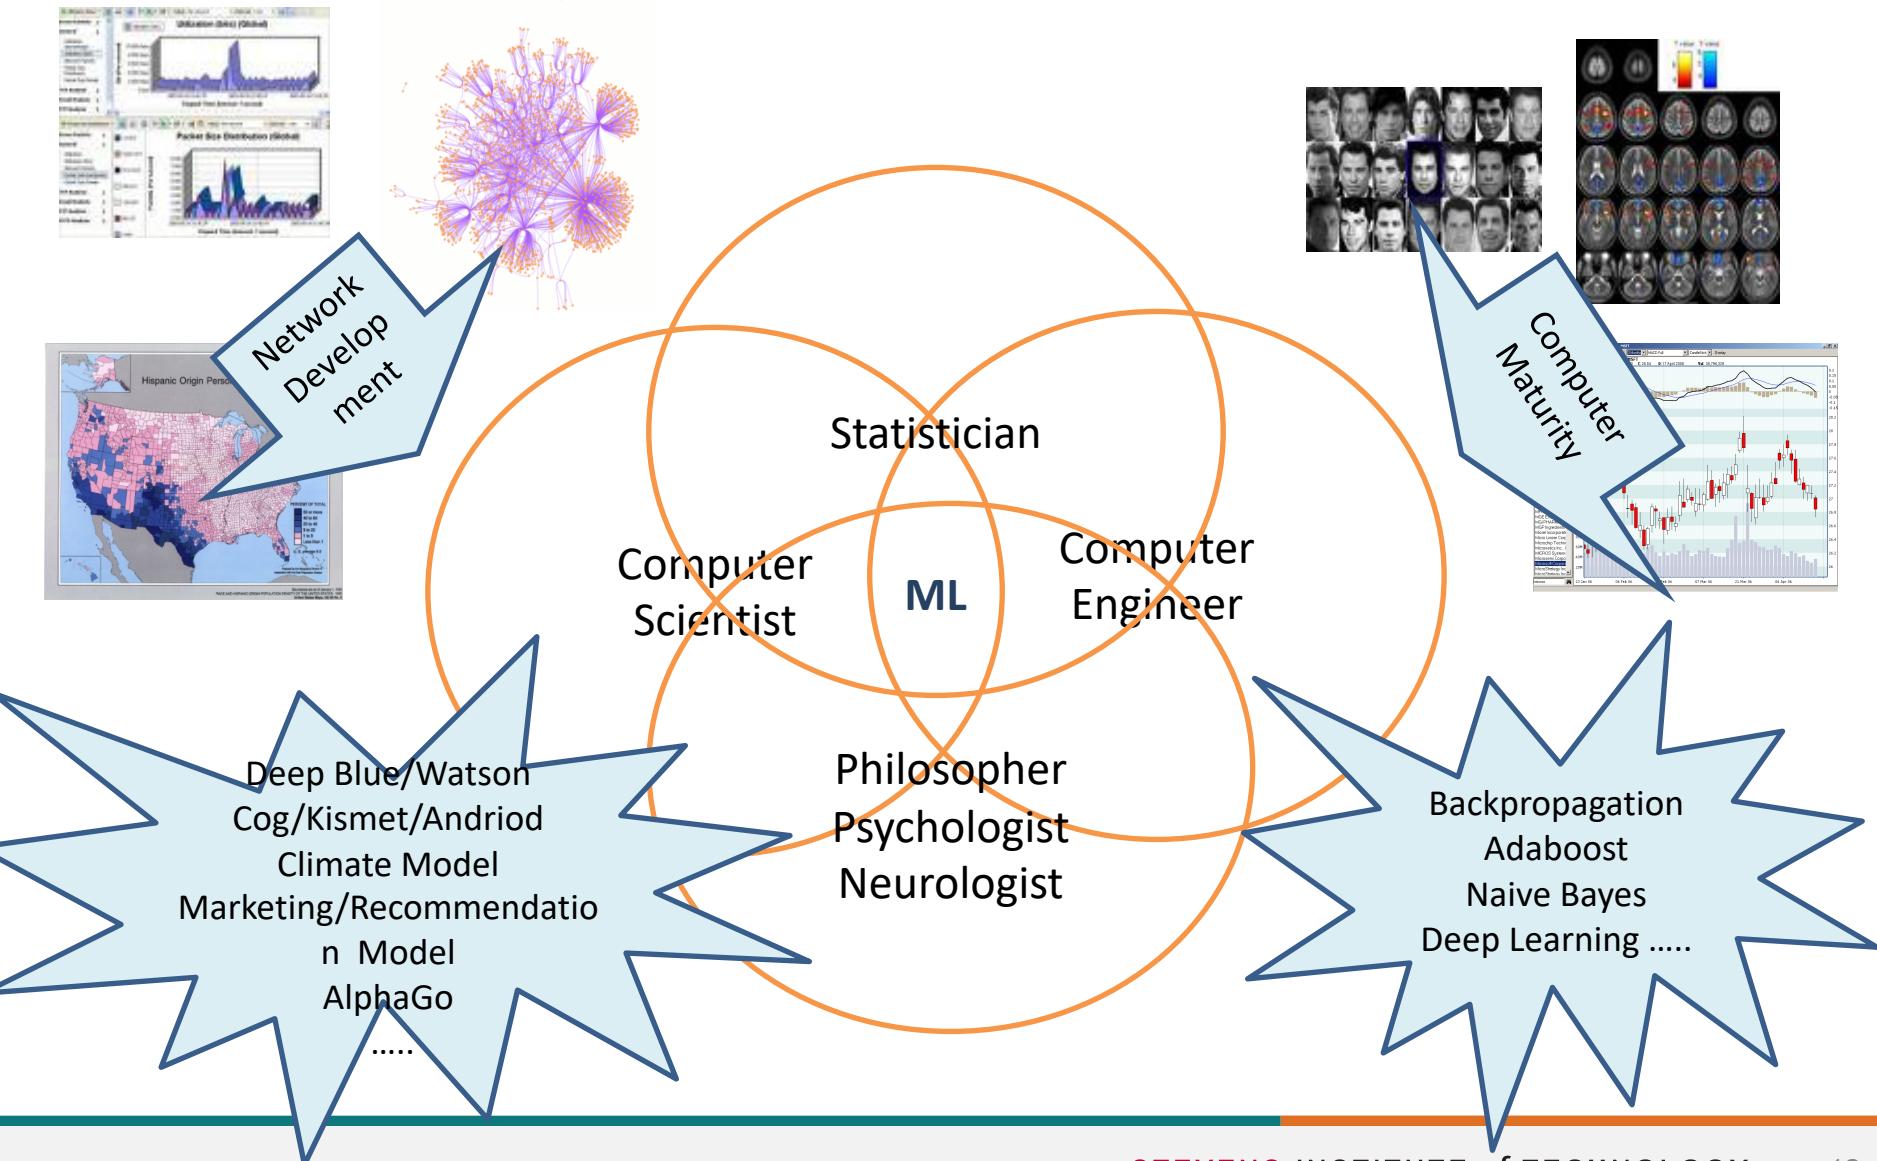


Exclusive-OR gate

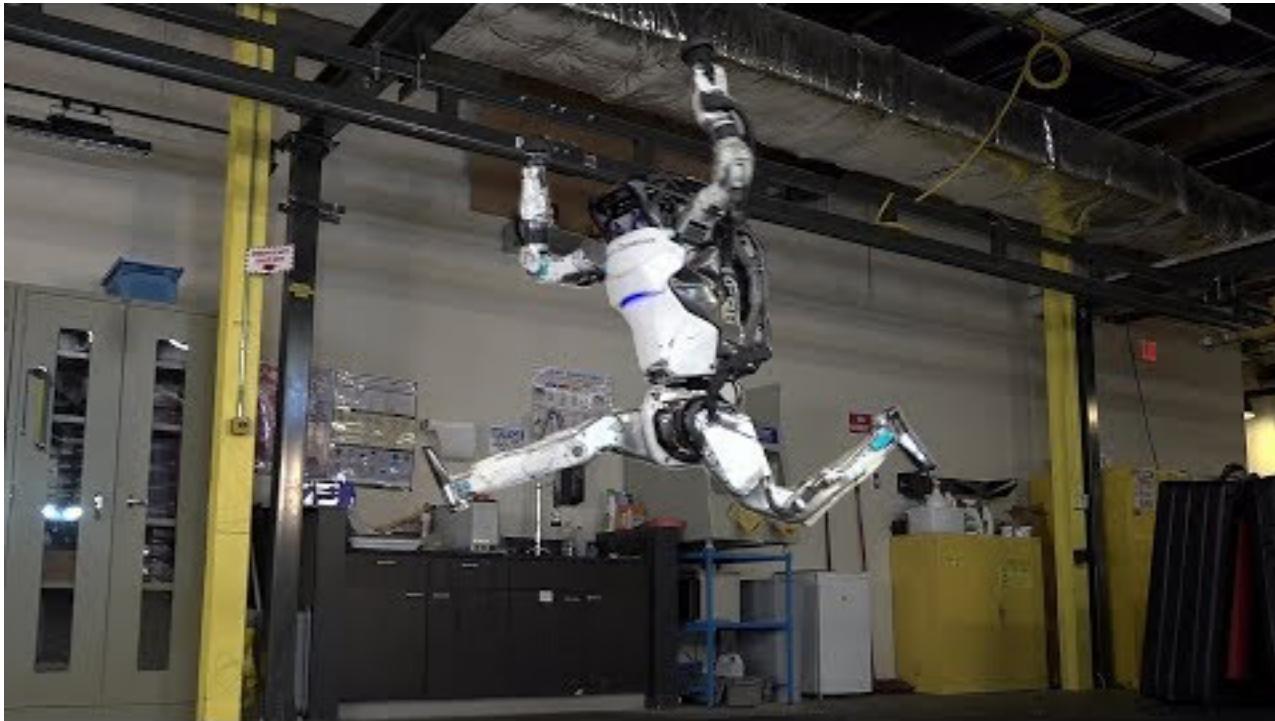


A	B	Output
0	0	0
0	1	1
1	0	1
1	1	0

After 1990's



Robots



Signal Processing

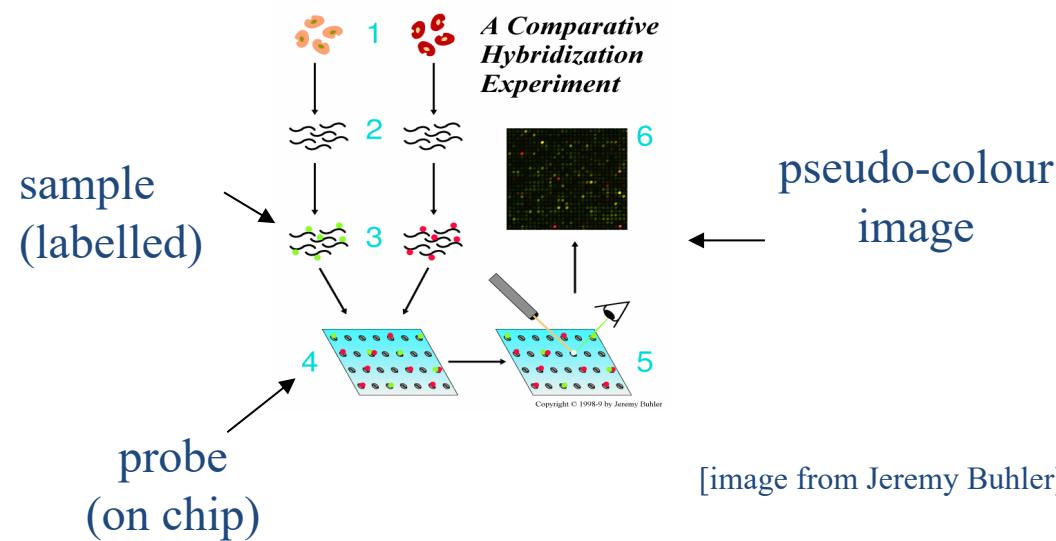
- Face Recognition



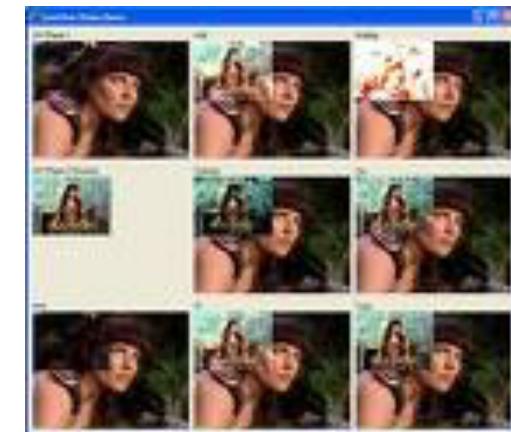
- Speech Recognition



- Microarray



- Video Processing





Business Analysis

Fraud Detection



Market Basket



Churn Prediction



Web Mining



Stock Market



Recommendation Systems



\$1 million, 10% improve,
4 years(2006 – 2009)

machine learning

Winnie Duan
Machine learning is fun

[Reply](#) [Forward](#)

FIAT 500 - Drive a New Italian Classic
Buy a FIAT 500 at a FIAT Studio.
[.com](#)

show details 11:06 PM (0 minutes ago) [Reply](#) [Forward](#)

Ads

Capital One Official Site
0% Intro rates. Great rewards.
No annual Fee. Apply now!
[www.CapitalOne.com](#)

Fly US - India @ \$729
Buy 1, Get Unlimited Tickets Free.
@ MakeMyTrip. Book Flights Now!
[us.makemytrip.com/Special-Fares](#)

American Express Business
The Right Card for Your Business.
Great Offers from American Express.
[www.AmericanExpress.com/Business](#)

More about...

- [Data Mining »](#)
- [Online Learning »](#)
- [Blended Learning Content »](#)
- [Learning Styles »](#)
- [Registration »](#)
- [Email Account »](#)
- [Hotmail Account »](#)
- [Account Software »](#)

Amazon

Recommend Link relates to “Machine Learning”

\$155.00 \$89.76
[View or edit your browsing history](#)

New For You:

Torn (Missing)
Margaret Peterson Haddix
Hardcover
\$16.99 **\$11.31**
Fix this recommendation

Percy Jackson pbk 5-book boxed set
Rick Riordan
Paperback
\$35.00 **\$23.10**
Fix this recommendation

[See more new releases](#)

\$99.95 \$74.69

The Death Cure (Maze Runner Trilogy)
James Dashner
Hardcover
\$17.99 **\$10.58**
Fix this recommendation

The Power of Six (Lorien Legacies)
Pittacus Lore
Hardcover
\$17.99 **\$8.98**
Fix this recommendation

\$89.95 \$57.91

Diary of a Wimpy Kid: Cabin Fever
Jeff Kinney
Hardcover
\$19.99 **\$7.67**
Fix this recommendation

The 39 Clues: Cahills vs. Vespers...
Gordon Korman
Hardcover
\$12.99 **\$8.11**
Fix this recommendation

\$5 billion, 25% of the annual sales

Climate/Healthcare Science

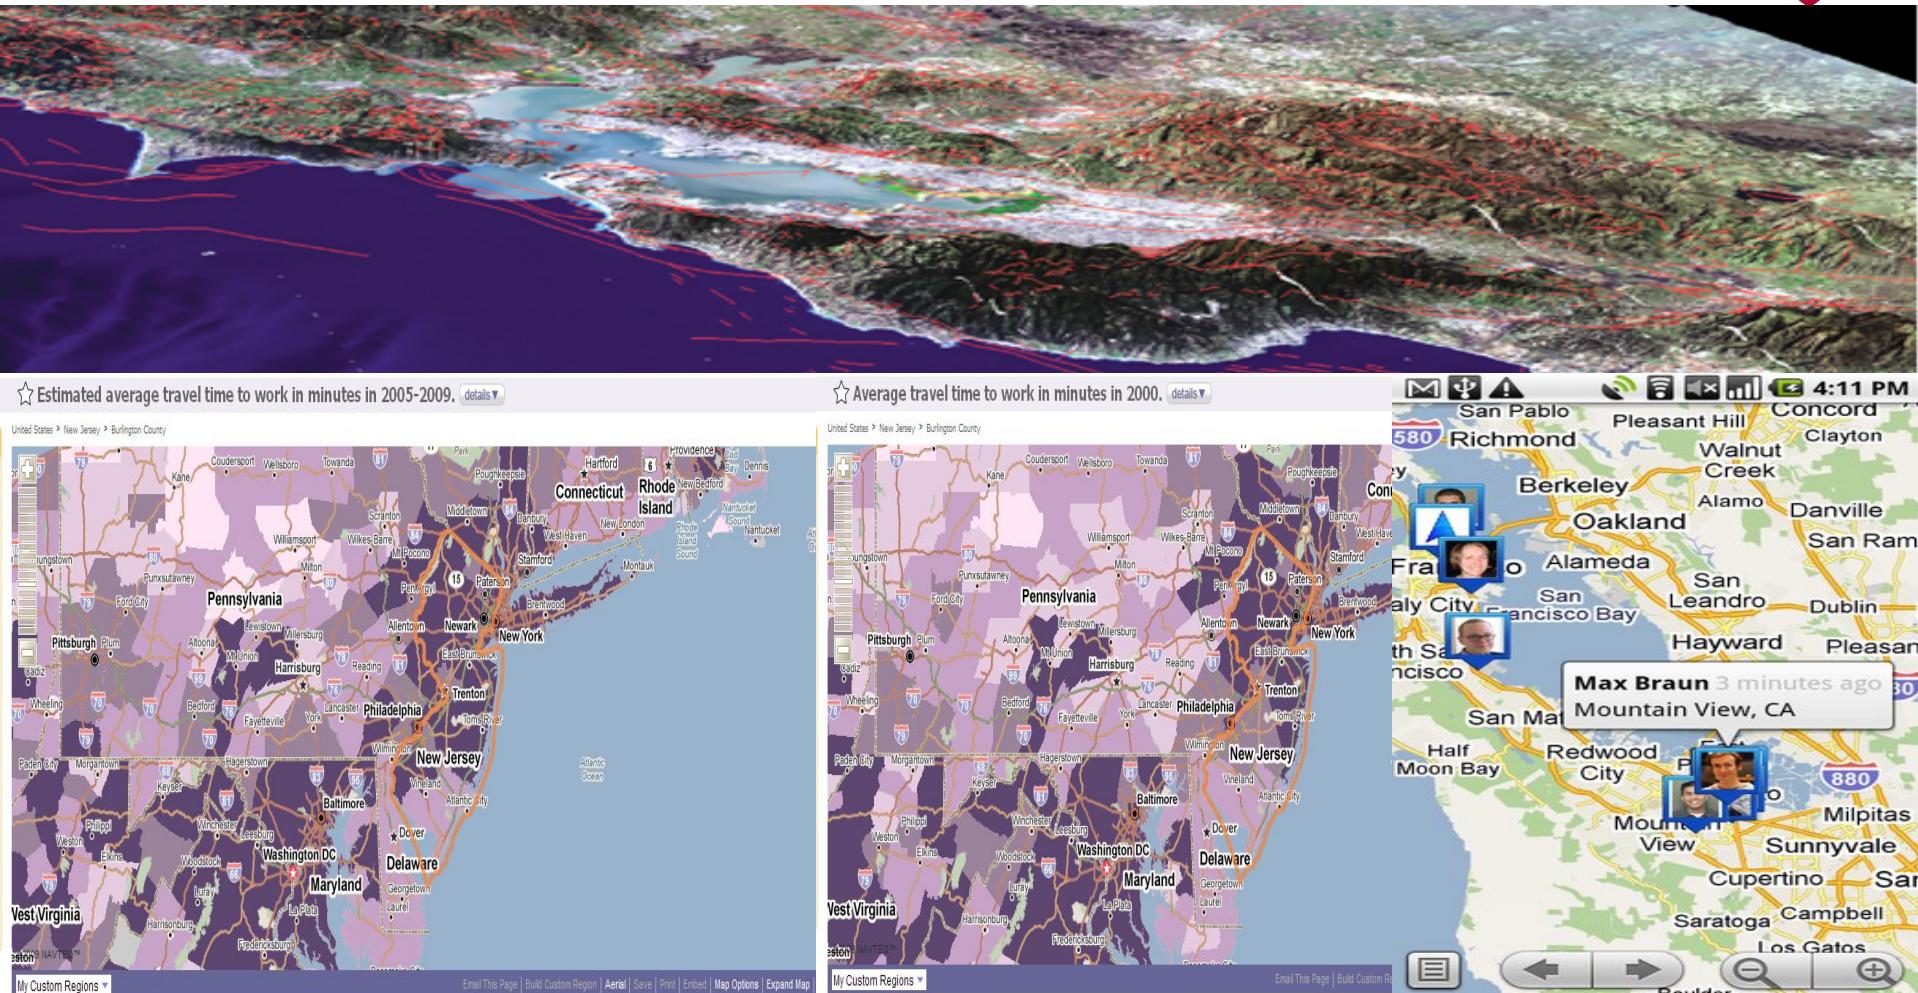
Healthcare



Climate Change



GIS and Location Science



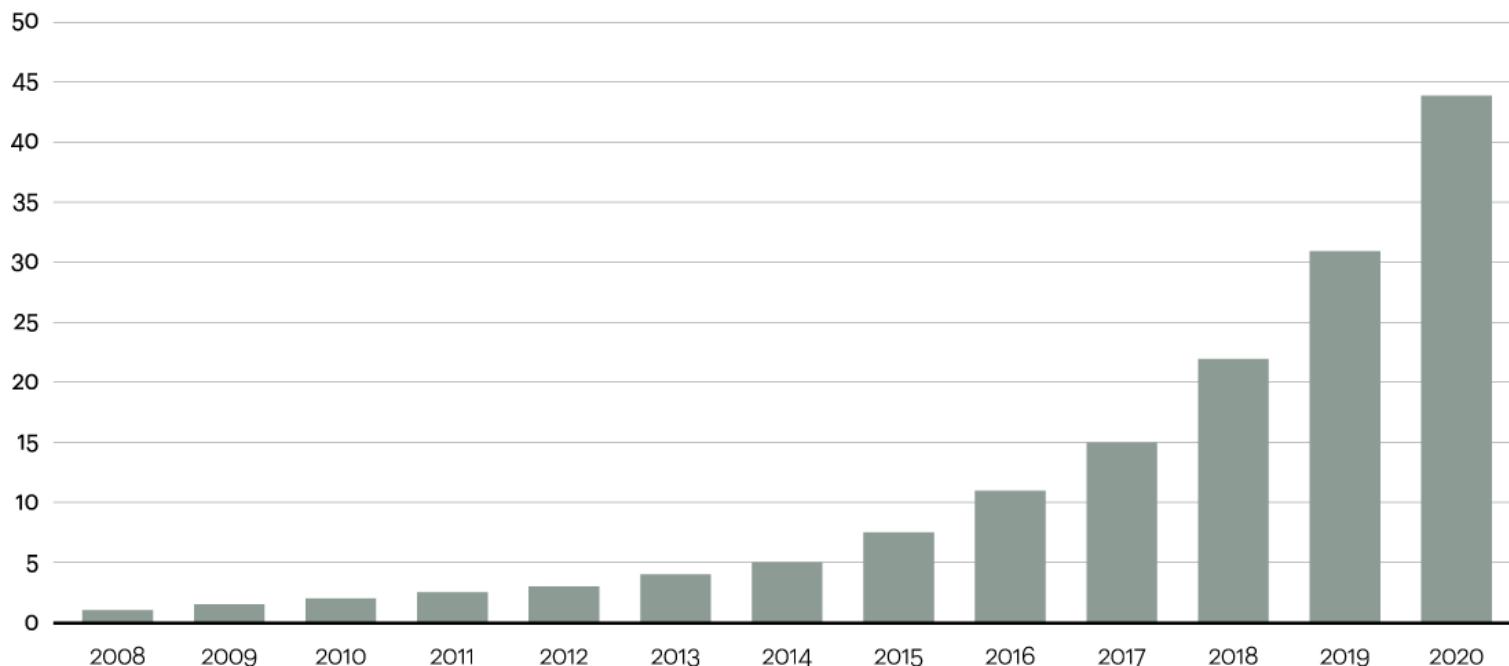
Complexity of Geographic Information System:
 High dimension Dynamic
 Spatial temporal Computation

Big Data Analytics

Figure 1

Data is growing at a 40 percent compound annual rate, reaching nearly 45 ZB by 2020

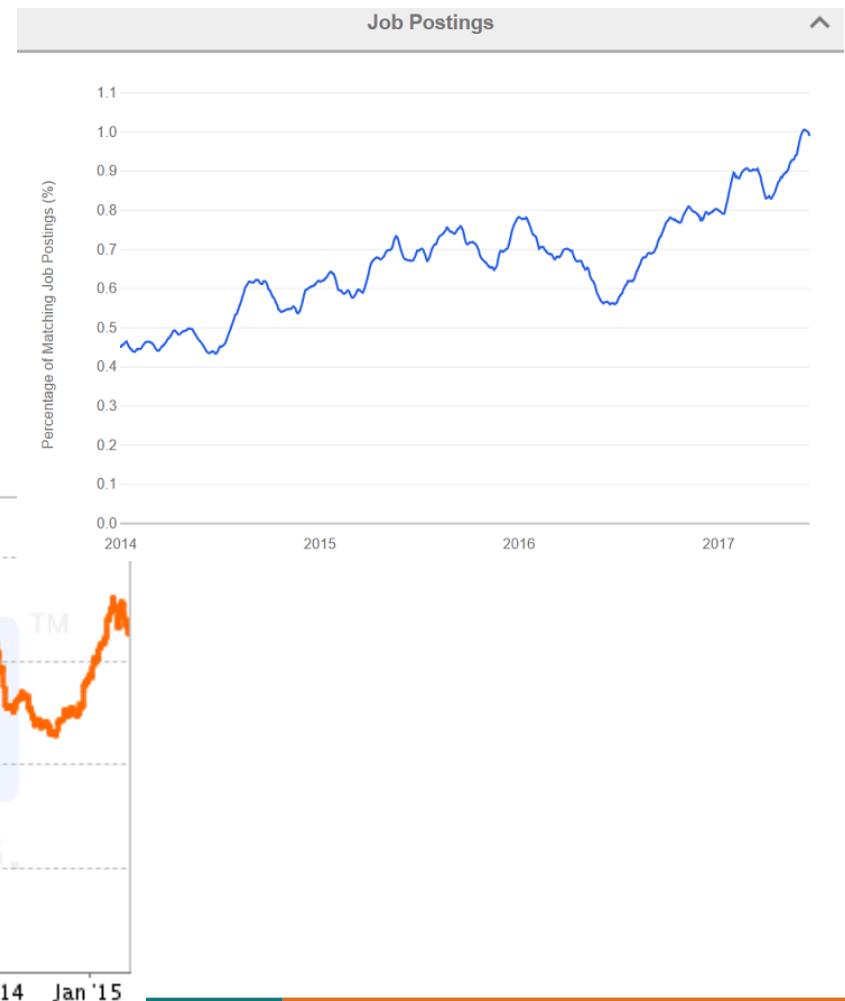
Data in zettabytes (ZB)



Source: Oracle, 2012

Machine Learning Jobs

“140,000 – 190,000 position shortage with deep expertise in statistics and machine learning” --- Mckinsey Big Data report





Challenges for Machine Learning

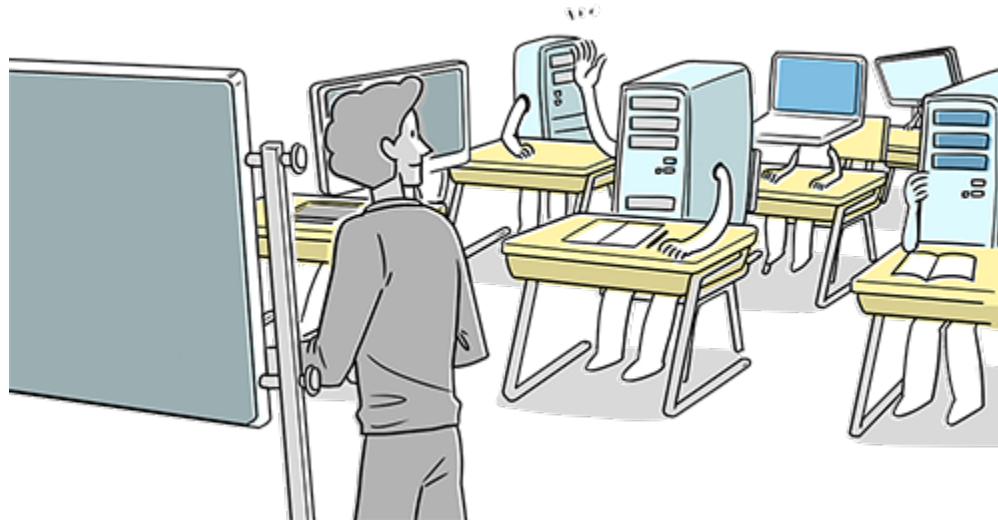
- Intelligent Systems
 - **Dynamic, incremental and adaptive learning algorithms;**
 - Limited understanding of intelligent behaviors;
 - **Lack of system level models and architectures to mimic brain-level intelligence;**
 - Embodiment of machine intelligence hardware within systems that learn through interaction with environment;
 - Systems design and implementation;
- Big Data
 - Platform/Storage/Index
 - Quality Uncertainty
 - **Efficient statistical/machine learning methods**
 - Intellectual property, liability
 - Visualization



Challenges for Machine Learning

- Theoretical foundations
 - Some results based on optimization and probability models
 - Lack rigorous mathematical formulation especially for recent deep learning algorithms
 - Verifiability of the system is another problem
- Security, privacy and ethical issues
 - Model reliability under attacks, e.g., pollution attacks
 - Data privacy, model privacy
 - Ethical use of machine learning in controllable ways

Again, What is Machine Learning?



Simple Answer

What is Machine Learning?

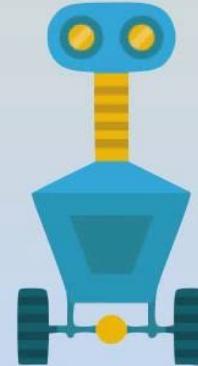
Learn from experience



Learn from experience



Follow instructions





Answer from Expert

“A computer program is said to learn from experience **E** with respect to some class of tasks **T** and performance measure **P**, if its performance at tasks in **T**, as measured by **P**, improves with experience **E**.“

--- Tom Mitchell



Why Learning?

Learning is used when:

- Human expertise does not exist (navigating on Mars),
- Humans are unable to explain their expertise (speech recognition)
- Solution changes in time
- Solution needs to be adapted to particular cases
- Data Volume are huge



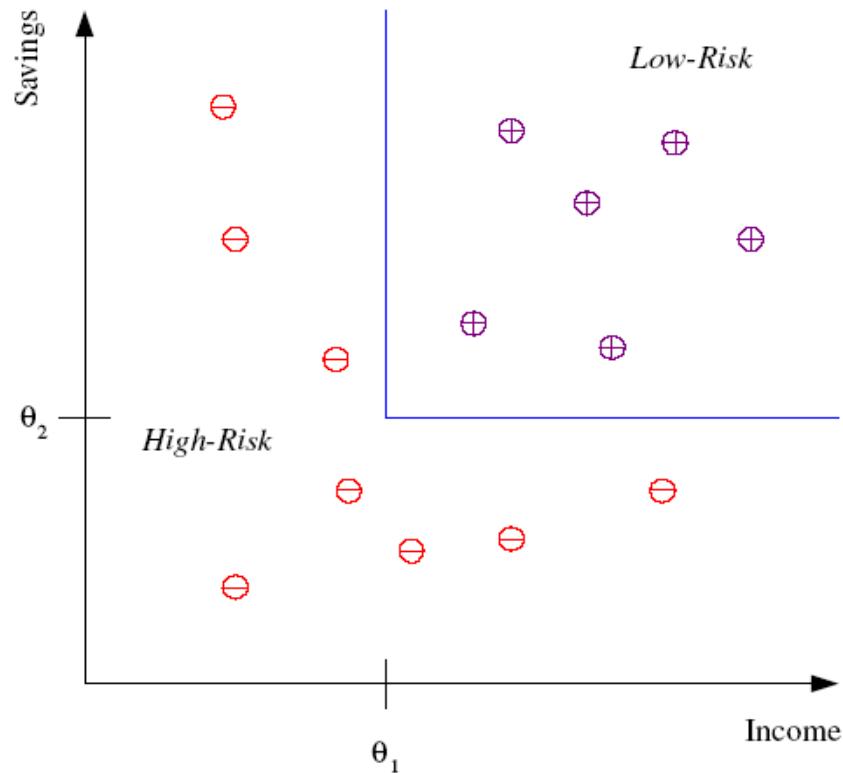
Learning Associations

- Basket analysis:
 $P(Y | X)$ probability that somebody who buys X also buys Y where X and Y are products/services.

Example: $P(\text{chips} | \text{beer}) = 0.7$

Classification

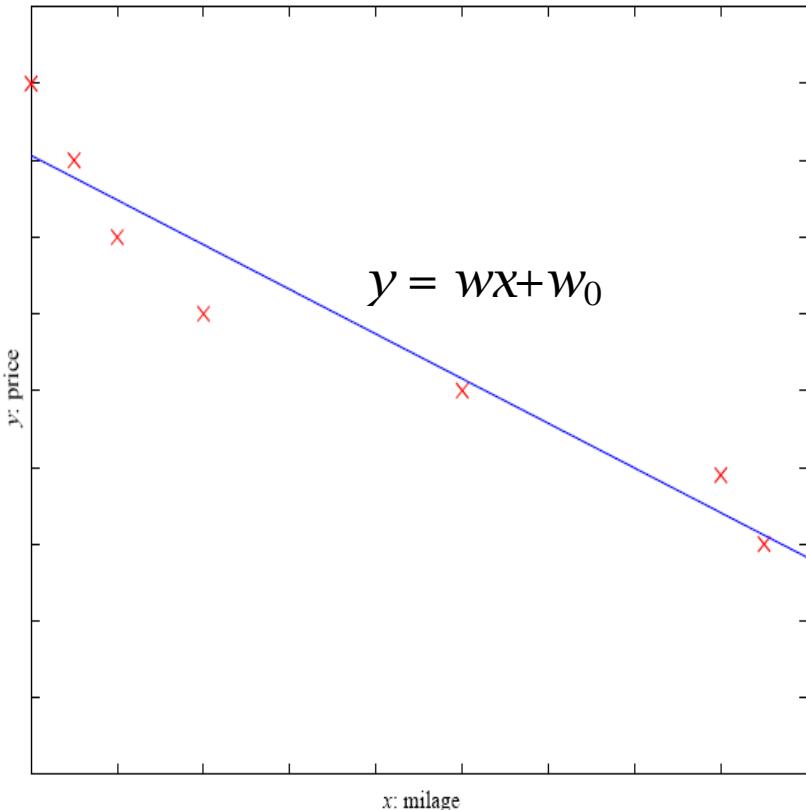
- Example: Credit scoring
- Differentiating between **low-risk** and **high-risk** customers from their *income* and *savings*



Discriminant: IF $income > \theta_1$ AND $savings > \theta_2$
 THEN **low-risk** ELSE **high-risk**

Regression

- Example: Price of used cars
- x : car attributes
 y : price
 $y = g(x | \theta)$
 $g(\cdot)$ model,
 θ parameters





Reinforcement Learning

- Learning a policy: A sequence of outputs
- No supervised output but delayed reward
- Credit assignment problem
- Game playing
- Robot in a maze
- Multiple agents, partial observability, ...



Categorization of Learning Systems

- Whether or not trained with human supervision
 - Supervised
 - k-NN, linear regression, logistic regression, support vector machines, decision trees and random forests, neural networks, etc.
 - Unsupervised
 - Clustering (e.g., k-means), dimensionality reduction (e.g., PCA), association rule learning, etc.
 - Semi-supervised
 - Reinforcement learning



Categorization of Learning Systems

- Whether or not learn incrementally on the fly
 - Online learning
 - Batch learning



Categorization of Learning Systems

- Whether compare new data to known data points, or build a predictive model
 - Instance-based learning
 - Model-based learning



Programming Tools

- Examples and assignments will use Python language. Jupyter notebook provides visualized interface.
- The following libraries are necessary for this course:
Scikit-learn (<http://scikit-learn.org>), Numpy, Matplotlib, Pandas, ...
- **Install Anaconda** (<https://www.anaconda.com/download/#macos>) first!
It includes the whole package including Python and related libraries.
- The examples and code of Aurelien Geron's textbook can be found at
<https://github.com/ageron/handson-ml>



Reference

The lecture notes in this lecture are based on the following textbooks:

T. M. Mitchell, Machine Learning, McGraw Hill, 1997. ISBN: 978-0-07-042807-2

C. M. Bishop, Pattern Recognition and Machine Learning, Springer, 2006, ISBN: 978-0-387-31073-2.

E. Alpaydin, Introduction to Machine Learning, MIT Press, 2004, ISBN 0-262-01211-1

Part of the slide materials were based on Dr. Rong Duan's Fall 2016 course CPE/EE 695A Applied Machine Learning at Stevens Institute of Technology.



STEVENS
INSTITUTE *of* TECHNOLOGY
THE INNOVATION UNIVERSITY®

stevens.edu

